

NAME : SK SHAREEF

REG NO: 21BCE9384

**Data Preprocessing **


```
#Import the Libraries.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
#Importing the dataset.
df=pd.read_csv("employee.csv")
```

```
df.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educati
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2

5 rows × 35 columns



```
df.shape
```

(1470, 35)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                        1470 non-null   object
18  MonthlyIncome                        1470 non-null   int64
19  MonthlyRate                           1470 non-null   int64
20  NumCompaniesWorked                   1470 non-null   int64
21  Over18                               1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                    1470 non-null   int64
24  PerformanceRating                    1470 non-null   int64
25  RelationshipSatisfaction              1470 non-null   int64
26  StandardHours                        1470 non-null   int64
27  StockOptionLevel                     1470 non-null   int64
28  TotalWorkingYears                    1470 non-null   int64
29  TrainingTimesLastYear                1470 non-null   int64
30  WorkLifeBalance                      1470 non-null   int64
31  YearsAtCompany                       1470 non-null   int64
32  YearsInCurrentRole                   1470 non-null   int64
```

9/26/23, 10:00 PM

Assignment 4 .ipynb - Colaboratory

33 YearsSinceLastPromotion1470 non-nullint64

34 YearsWithCurrManager1470 non-nullint64

dtypes: int64(26), object(9)

memory usage: 402.1+ KB

df.describe()

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	1470.000000	1470.000000
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	2.721769	65.891156
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	1.093082	20.329428
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	1.000000	30.000000
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	2.000000	48.000000
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	3.000000	66.000000
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	4.000000	83.750000
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	4.000000	100.000000

8 rows × 26 columns

#Checking for Null Values.

df.isnull().any()

AgeFalse

AttritionFalse

BusinessTravelFalse

DailyRateFalse

DepartmentFalse

DistanceFromHomeFalse

EducationFalse

EducationFieldFalse

EmployeeCountFalse

EmployeeNumberFalse

EnvironmentSatisfactionFalse

GenderFalse

HourlyRateFalse

JobInvolvementFalse

JobLevelFalse

JobRoleFalse

JobSatisfactionFalse

MaritalStatusFalse

MonthlyIncomeFalse

MonthlyRateFalse

NumCompaniesWorkedFalse

Over18False

OverTimeFalse

PercentSalaryHikeFalse

PerformanceRatingFalse

RelationshipSatisfactionFalse

StandardHoursFalse

StockOptionLevelFalse

TotalWorkingYearsFalse

TrainingTimesLastYearFalse

WorkLifeBalanceFalse

YearsAtCompanyFalse

YearsInCurrentRoleFalse

YearsSinceLastPromotionFalse

YearsWithCurrManagerFalse

dtype: bool

df.isnull().sum()

Age0

Attrition0

BusinessTravel0

DailyRate0

Department0

DistanceFromHome0

Education0

EducationField0

EmployeeCount0

EmployeeNumber0

EnvironmentSatisfaction0

Gender0

HourlyRate0

JobInvolvement0

JobLevel0

JobRole0


JobSatisfaction0

MaritalStatus0

```
MonthlyIncome      0
MonthlyRate        0
NumCompaniesWorked  0
Over18             0
OverTime           0
PercentSalaryHike   0
PerformanceRating   0
RelationshipSatisfaction  0
StandardHours       0
StockOptionLevel    0
TotalWorkingYears   0
TrainingTimesLastYear  0
WorkLifeBalance     0
YearsAtCompany      0
YearsInCurrentRole   0
YearsSinceLastPromotion  0
YearsWithCurrManager  0
dtype: int64
```

DATA VISUALISATION

```
#histogram of Age
plt.hist(df['Age'], bins=20, color='skyblue', edgecolor='black')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.title('Age Distribution')
plt.show()
```




```
# countplot
sns.countplot(x='Department', data=df, palette='pastel')
plt.xlabel('Department')
plt.ylabel('Count')
plt.title('Number of Employees in Each Department')
plt.show()
```




OUTLIERS DETECTION


```
#checking for age column  
sns.boxplot(df.Age)
```



```
#Checking for Daily rate  
sns.boxplot(df.DailyRate)
```




```
#checking for DailyRate
sns.boxplot(df.DailyRate)
```



```
df.head(3)
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumb
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	

3 rows × 35 columns



```
#checking for standardhours
sns.boxplot(df.StandardHours)
```

```
<Axes: >
0.1 |-----|
No outliers found so we continue with next step
0.5 |-----|
ENCODING
|-----|
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
0.5 |-----|
df.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns

```
#performing label encoding some satisfied columns
columns = ['BusinessTravel', 'Department', 'Gender', 'Over18', 'OverTime']
df[columns]=df[columns].apply(le.fit_transform)

#performing one hot encoding on some satisfied columns
one_hot_columns = ['JobRole', 'EducationField', 'MaritalStatus']
df = pd.get_dummies(df, columns=one_hot_columns, drop_first=True)

df.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSa
0	41	Yes	2	1102	2	1	2	1	1	
1	49	No	1	279	1	8	1	1	2	
2	37	Yes	2	1373	1	2	2	1	4	
3	33	No	1	1392	1	3	4	1	5	
4	27	No	2	591	1	2	1	1	7	

5 rows × 47 columns

```
df.columns

Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
      'DistanceFromHome', 'Education', 'EmployeeCount', 'EmployeeNumber',
      'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
      'JobLevel', 'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate',
      'NumCompaniesWorked', 'Over18', 'OverTime', 'PercentSalaryHike',
      'PerformanceRating', 'RelationshipSatisfaction', 'StandardHours',
      'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
      'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
      'YearsSinceLastPromotion', 'YearsWithCurrManager',
      'JobRole_Human Resources', 'JobRole_Laboratory Technician',
      'JobRole_Manager', 'JobRole_Manufacturing Director',
      'JobRole_Research Director', 'JobRole_Research Scientist',
      'JobRole_Sales Executive', 'JobRole_Sales Representative',
      'EducationField_Life Sciences', 'EducationField_Marketing',
      'EducationField_Medical', 'EducationField_Other',
      'EducationField_Technical Degree', 'MaritalStatus_Married',
      'MaritalStatus_Single'],
      dtype='object')
```

SPLITTING IN TO DEPENDENT AND INDEPENDENT

```
# Independent variables (features)
x = df[['Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EmployeeCount', 'EmployeeNumber',
        'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel', 'JobSatisfaction',
        'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'PercentSalaryHike',
        'PerformanceRating', 'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
        'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany',
        'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager',
        'BusinessTravel', 'Department', 'Gender', 'Over18', 'OverTime',
        'JobRole_Human Resources', 'JobRole_Laboratory Technician', 'JobRole_Manager',
        'JobRole_Manufacturing Director', 'JobRole_Research Director', 'JobRole_Research Scientist',
        'JobRole_Sales Executive', 'JobRole_Sales Representative',
        'EducationField_Life Sciences', 'EducationField_Marketing', 'EducationField_Medical',
        'EducationField_Other', 'EducationField_Technical Degree',
        'MaritalStatus_Married', 'MaritalStatus_Single']]

# Dependent variable (target)
y = df['Attrition']
```

FEATURE SACLING

```
#feature scaling
from sklearn.preprocessing import MinMaxScaler
ms=MinMaxScaler()
x_scaled=pd.DataFrame(ms.fit_transform(x),columns=x.columns)
```

SPLITTING DATA INTO TRAIN AND SET

```
#TRAIN TEST AND SPLIT
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x_scaled,y,test_size =0.2,random_state =42)

x_train.shape,x_test.shape,y_train.shape,y_test.shape

((1176, 46), (294, 46), (1176,), (294,))
```

MODEL BUILDING

- o Import the model building Libraries
- o Initializing the model
- o Training and testing the model
- o Evaluation of Model
- o Save the Model

1.LOGISTIC REGRESION

2. DECISION TREE
3. RANDOM FOREST

```
#LOGISTIC REGRESSION
from sklearn.linear_model import LogisticRegression
L_model=LogisticRegression()
```

```
L_model.fit(x_train,y_train)
```

```
▼ LogisticRegression
LogisticRegression()
```

```
pred_L=L_model.predict(x_test)
```

EVALUATION OF THE CLASSIFICATION MODEL

```
#Accuracy score
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,roc_auc_score,roc_curve

accuracy_score(y_test,pred)
```

0.8877551020408163

```
confusion_matrix(y_test,pred)

array([[248,  7],
       [ 26, 13]])
```

```
pd.crosstab(y_test,pred)
```

col_0	No	Yes
Attrition		
No	248	7
Yes	26	13

DECISION TREE MODEL

```
from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
```

```
dtc.fit(x_train,y_train)
```

▼ DecisionTreeClassifier

DecisionTreeClassifier()

```
D_pred=dtc.predict(x_test)
```

Evaluation of classification model

```
#Accuracy score
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,roc_auc_score,roc_curve
```

```
accuracy_score(y_test,D_pred)

0.7789115646258503
```

```
confusion_matrix(y_test,D_pred)

array([[222, 33],
       [ 32,  7]])
```

```
pd.crosstab(y_test,D_pred)
```

col_0	No	Yes
Attrition		
No	222	33
Yes	32	7

```
print(classification_report(y_test,D_pred))
```

	precision	recall	f1-score	support
No	0.87	0.87	0.87	255
Yes	0.17	0.18	0.18	39
accuracy			0.78	294
macro avg	0.52	0.53	0.52	294
weighted avg	0.78	0.78	0.78	294

HYPER PARAMETER TUNING

```
from sklearn import tree
plt.figure(figsize=(25,15))
tree.plot_tree(dtc,filled=True)
```



```

from sklearn.model_selection import GridSearchCV
parameter={
    'criterion':['gini','entropy'],
    'splitter':['best','random'],
    'max_depth':[1,2,3,4,5],
    'max_features':['auto', 'sqrt', 'log2']
}

grid_search=GridSearchCV(estimator=dtc,param_grid=parameter,cv=5,scoring="accuracy")

grid_search.fit(x_train,y_train)

grid_search.best_params_

{'criterion': 'gini',
 'max_depth': 5,
 'max_features': 'sqrt',
 'splitter': 'best'}

dtc_cv=DecisionTreeClassifier(criterion= 'gini',
                              max_depth=5,
                              max_features='sqrt',
                              splitter='best')
dtc_cv.fit(x_train,y_train)

```

```

▼ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5, max_features='sqrt')

```

```
pred=dtc_cv.predict(x_test)
```

```
print(classification_report(y_test,D_pred))
```

	precision	recall	f1-score	support
No	0.87	0.87	0.87	255
Yes	0.17	0.18	0.18	39
accuracy			0.78	294
macro avg	0.52	0.53	0.52	294
weighted avg	0.78	0.78	0.78	294

RANDOM FOREST

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report

# Initialize the Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the classifier on the training data
rf_classifier.fit(x_train, y_train)

# Predict on the test data
R_pred = rf_classifier.predict(x_test)

# Evaluate the model
accuracy = accuracy_score(y_test, R_pred)
report = classification_report(y_test, R_pred)

print(f"Accuracy: {accuracy}")
print("\nClassification Report:")
print(report)

```

Accuracy: 0.8775510204081632

Classification Report:

	precision	recall	f1-score	support
No	0.88	1.00	0.93	255
Yes	0.80	0.10	0.18	39
accuracy			0.88	294
macro avg	0.84	0.55	0.56	294

weighted avg	0.87	0.88	0.83	294
--------------	------	------	------	-----